

Data Analysis With Apache Flink



Aljoscha Krettek / Till Rohrmann

Flink committers

Co-founders @ *data Artisans*

aljoscha@apache.org / trohrmann@apache.org

What is Apache Flink?



Functional
API

Relational
API

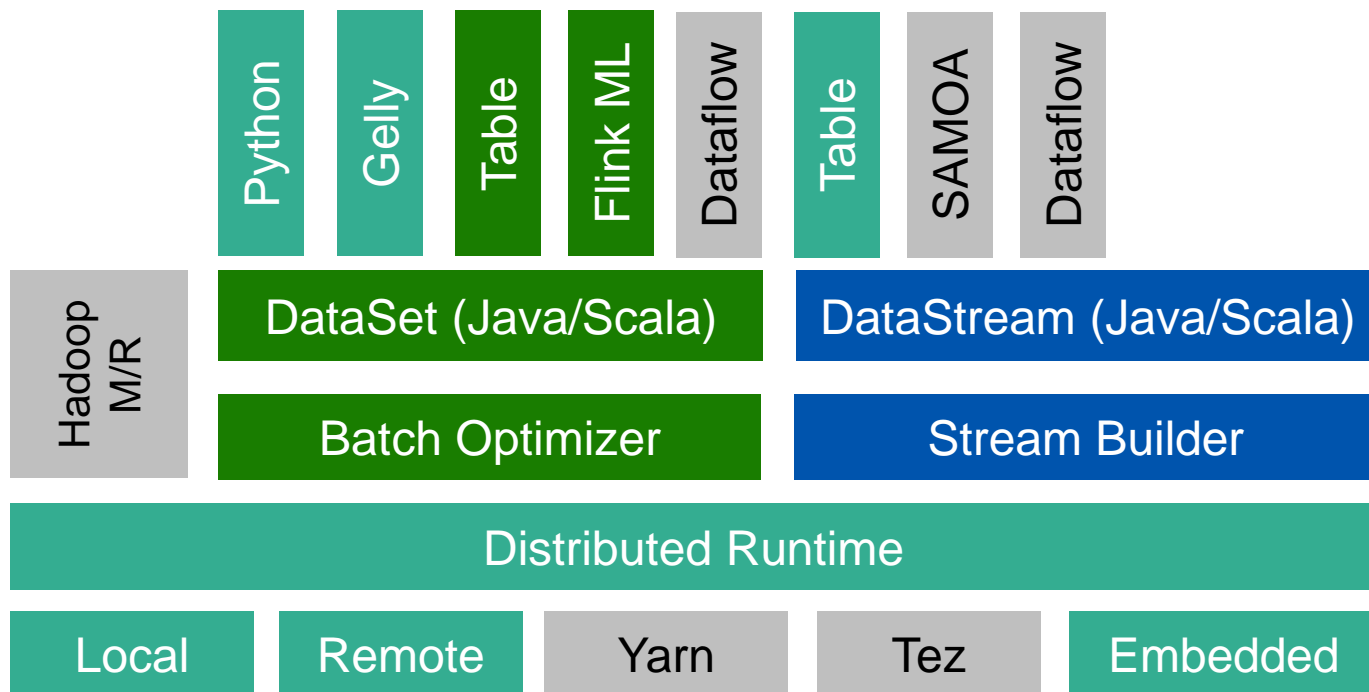
Graph API

Machine
Learning

...

Iterative Dataflow Engine

Apache Flink Stack



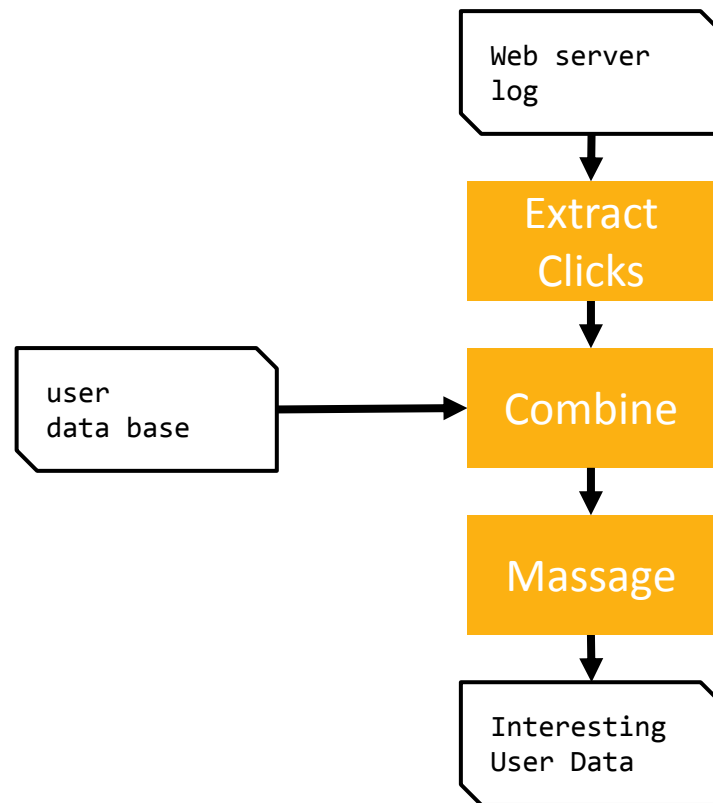
**current Flink master + few PRs*

Example Use Case: Log Analysis

What Seems to be the Problem?



- Collect clicks from a webserver log
- Find interesting URLs
- Combine with user data



The Execution Environment



- Entry point for all Flink programs
- Creates DataSets from data sources

```
ExecutionEnvironment env =  
    ExecutionEnvironment.getExecutionEnvironment();
```

Getting at Those Clicks



```
post /foo/bar... 313  
get /data/pic.jpg 128  
post /bar/baz... 128  
post /hello/there... 42
```

```
DataSet<String> log = env.readTextFile("hdfs:///log");
```

```
DataSet<Tuple2<String, Integer>> clicks = log.flatMap(
```

```
(String line, Collector<Tuple2<String, Integer>> out) ->  
  String[] parts = in.split("*magic regex*");  
  if (isClick(parts)) {  
    out.collect(new Tuple2<>(parts[1], Integer.parseInt(parts[2])));  
  }  
}  
)
```

The Table Environment



- Environment for dealing with Tables
- Converts between DataSet and Table

```
TableEnvironment tableEnv = new TableEnvironment();
```


Counting those Clicks



```
Table clicksTable = tableEnv.toTable(clicks, "url, userId");
```

```
Table urlClickCounts = clicksTable  
  .groupBy("url, userId")  
  .select("url, userId, url.count as count");
```

Getting the User Information



```
Table userInfo = tableEnv.toTable(..., "name, id, ...");
```

```
Table resultTable = urlClickCounts.join(userInfo)  
  .where("userId = id && count > 10")  
  .select("url, count, name, ...");
```

The Final Step



```
class Result {  
    public String url;  
    public int count;  
    public String name;  
    ...  
}
```

```
DataSet<Result> set =  
    tableEnv.toSet(resultTable, Result.class);
```

```
DataSet<Result> result =  
    set.groupBy("url").reduceGroup(new ComplexOperation());
```

```
result.writeAsText("hdfs:///result");  
env.execute();
```

What happens under the hood?

From Program to Dataflow



```
ExecutionEnvironment env = TableEnvironment tableEnv = new TableEnvironment();
TableEnvironment tableEnv = new TableEnvironment();

DataSet<String> log = env.readTextFile("hdfs:///log");

DataSet<Tuple2<String, Integer>> clicks = log.flatMap(
    new FlatMapFunction<String, Tuple2<String, Integer>>() {
        public void flatMap(String in, Collector<Tuple2<String, Integer>> out) {
            String[] parts = in.split("magic regex");
            if (parts[0].equals("click")) {
                out.collect(new Tuple2<>(parts[1], Integer.parseInt(parts[4])));
            }
        }
    }
);

Table clicksTable = tableEnv.toTable(clicks, "url, userId");

Table urlClickCounts = clicksTable
    .groupBy("url, userId")
    .select("url, userId, count as count");

Table userInfo = tableEnv.toTable(..., "name, id, ...");


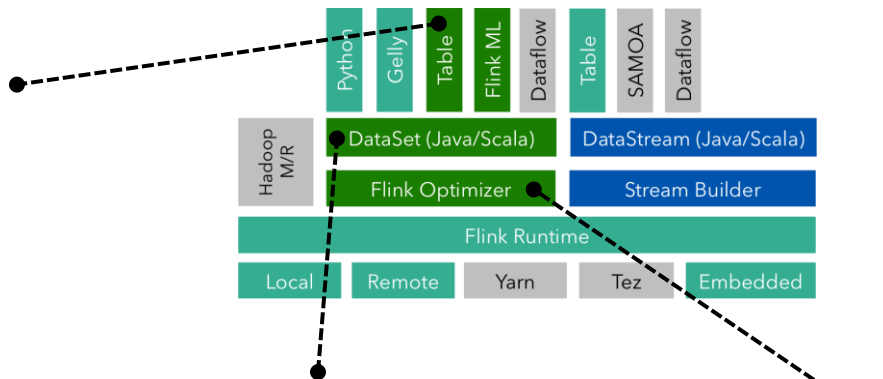

Table resultTable = urlClickCounts.join(userInfo)
    .where("userId = id && count > 10")
    .select("url, count, name, ...");

DataSet<Result> result = tableEnv.toSet(resultTable, Result.class);


result.writeAsText("hdfs:///result");

env.execute();
```

Flink Program




Dataflow Plan



Optimized Plan

Distributed Execution




Advanced Analysis: Website Recommendation

Going Further



- Log analysis result:
Which user visited how often which web site
- Which other websites might they like?
- Recommendation by collaborative filtering



Collaborative Filtering















- Recommend items based on users with similar preferences
- Latent factor models capture underlying characteristics of items and preferences of user
- Predicted preference: $\hat{r}_{u,i} = \mathbf{x}_u^T \mathbf{y}_i$

Matrix Factorization



$$R \gg X^T Y$$

				
	5	0	5	
	4	2	0	
	0	3	3	
				  
	a	1.1	0.8	0.7
	b	1.6	1.1	1
	c	0.7	0.5	0.4


				
a	1	0.6	1	
b	1.5	0.9	1.4	
c	0.6	0.4	0.6	

$$\min_{X,Y} \sum_{u,i} (r_{u,i} - x_u^T y_i)^2 + \lambda \sum_u \|x_u\|^2 + \lambda \sum_i \|y_i\|^2$$

Alternating least squares



- Iterative approximation
 1. Fix X and optimize Y
 2. Fix Y and optimize X
- Communication and computation intensive



Matrix Factorization Pipeline




```
val featureExtractor = HashingFT()
val factorizer = ALS()

val pipeline = featureExtractor.chain(factorizer)

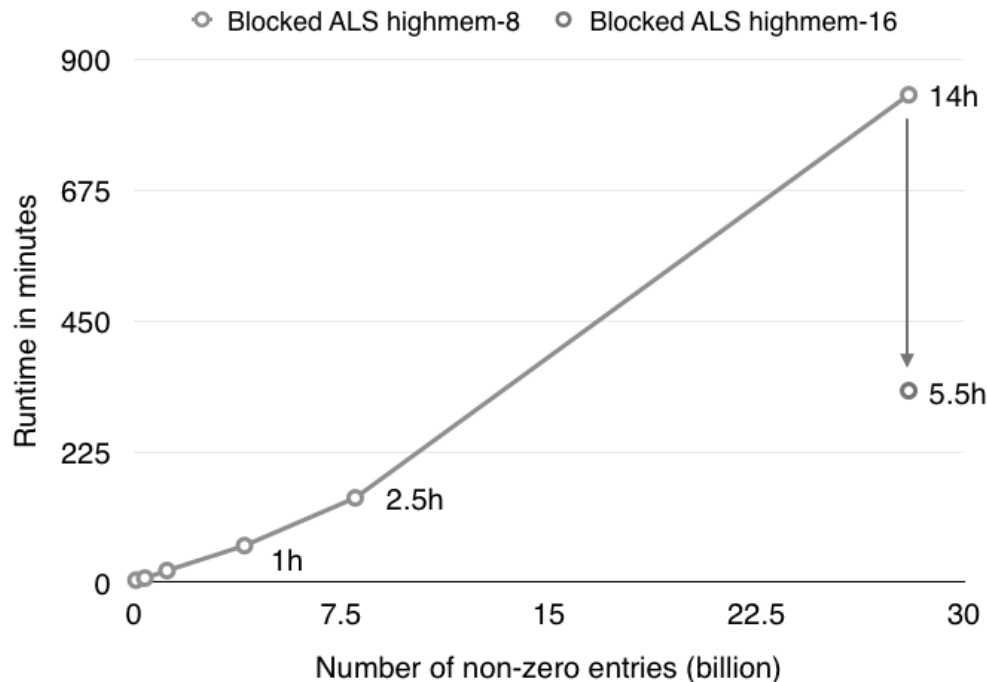
val clickstreamDS =
  env.readCsvFile[(String, String, Int)](clickStreamData)

val parameters = ParameterMap()
  .add(HashingFT.NumFeatures, 1000000)
  .add(ALS.Iterations, 10)
  .add(ALS.NumFactors, 50)
  .add(ALS.Lambda, 1.5)

val factorization = pipeline.fit(clickstreamDS, parameters)
```



Does it Scale?



Scale of Netflix or Spotify





- 40 node GCE cluster, highmem-8
- 10 ALS iteration with 50 latent factors
- Based on Spark MLlib's implementation

What Else Can You Do?



- Classification using SVMs
 - Conversion goal prediction
- Clustering
 - Visitor segmentation
- Multiple linear regression
 - Visitor prediction



Closing

What Have You Seen?



- Flink is a general-purpose analytics system
- Highly expressive Table API
- Advanced analysis with Flink's machine learning library
- Jobs are executed on powerful distributed dataflow engine

Flink Roadmap for 2015



- Additions to Machine Learning library
- Streaming Machine Learning
- Support for interactive programs
- Optimization for Table API queries
- SQL on top of Table API



— 2015 —

Flink *Forward*

— BERLIN 12/13 OCT —



flink.apache.org
@ApacheFlink

Backup Slides

WordCount in DataSet API



```
case class Word (word: String, frequency: Int)

val env = ExecutionEnvironment.getExecutionEnvironment()

val lines = env.readTextFile(...)

lines
  .flatMap {line => line.split(" ").map(word => Word(word,1))}
  .groupBy("word").sum("frequency")
  .print()

env.execute()
```

Java and Scala APIs offer the same functionality.

Log Analysis Code



```
ExecutionEnvironment env = TableEnvironment tableEnv = new TableEnvironment();
TableEnvironment tableEnv = new TableEnvironment();
```

```
DataSet<String> log = env.readTextFile("hdfs:///log");
```

```
DataSet<Tuple2<String, Integer>> clicks = log.flatMap(
    new FlatMapFunction<String, Tuple2<String, Integer>>() {

        public void flatMap(String in, Collector<Tuple2<>> out) {
            String[] parts = in.split("**magic regex*");
            if (parts[0].equals("click")) {
                out.collect(new Tuple2<>(parts[1], Integer.parseInt(parts[4])));
            }
        }
    });
```

```
Table clicksTable = tableEnv.toTable(clicks, "url, userId");
```

```
Table urlClickCounts = clicksTable
    .groupBy("url, userId")
    .select("url, userId, url.count as count");
```

```
Table userInfo = tableEnv.toTable(..., "name, id, ...");
```

```
Table resultTable = urlClickCounts.join(userInfo)
    .where("userId = id && count > 10")
    .select("url, count, name, ...");
```

```
DataSet<Result> result = tableEnv.toSet(resultTable, Result.class);
```

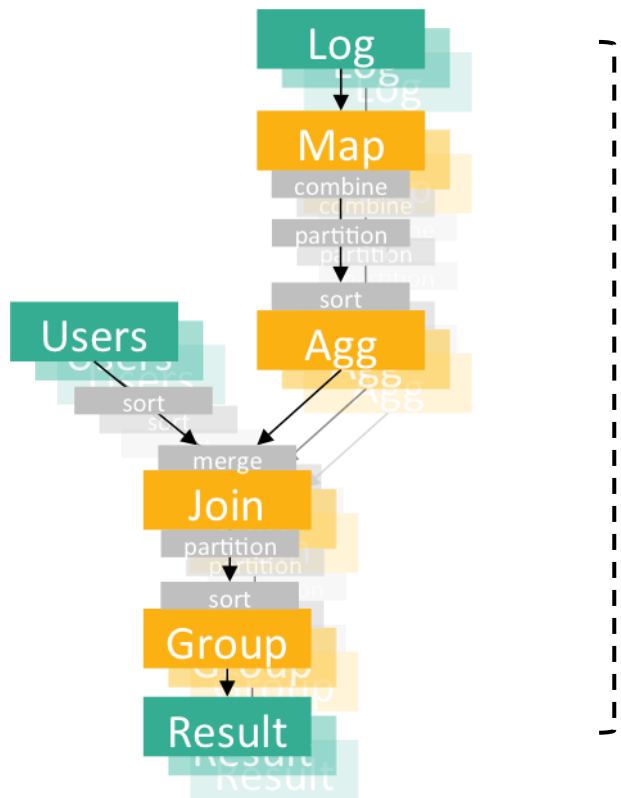
```
result.writeAsText("hdfs:///result");
```

```
env.execute();
```

Log Analysis Dataflow Graph



Pipelined Execution



Note: Intermediate DataSets are not necessarily “created”!

Only 1 Stage
(depending on join strategy)



Data transfer in-memory
and disk if needed

API in a Nutshell



- Element-wise
 - `map`, `flatMap`, `filter`
- Group-wise
 - `groupBy`, `reduce`, `reduceGroup`, `combineGroup`, `mapPartition`, `aggregate`, `distinct`
- Binary
 - `join`, `coGroup`, `union`, `cross`
- Iterations
 - `iterate`, `iterateDelta`
- Physical re-organization
 - `rebalance`, `partitionByHash`, `sortPartition`
- Streaming
 - `window`, `windowMap`, `coMap`, ...