

Apache Flink



Stephan Ewen

Flink committer


co-founder / CTO @ *data Artisans*

@StephanEwen

What is Apache Flink?



A "use-case complete" framework to unify
batch & stream processing



What is Apache Flink?



An engine that puts equal emphasis on streaming and batch processing

Real-time data streams





Event logs

Kafka, RabbitMQ, ...

Historic data

HDFS, JDBC, ...



ETL, Graphs,
Machine Learning
Relational, ...

Low latency
windowing,
aggregations, ...

What is Apache Flink?



- Large-scale data processing engine
- Easy and powerful APIs for *batch and real-time streaming* analysis (Java / Scala)
- Backed by a robust execution backend
 - with true streaming capabilities,
 - sophisticated windowing mechanisms,
 - custom memory manager,
 - native iteration execution,
 - and a cost-based optimizer.

Cornerpoints of Flink Design



Flexible Data Streaming Engine

- *Low Latency Stream Proc.*
- *Highly flexible windows*

Robust Algorithms on Managed Memory

- *No OutOfMemory Errors*
- *Scales to very large JVMs*
- *Efficient and robust processing*

High-level APIs, beyond key/value pairs

- *Java/Scala/Python (upcoming)*
- *Relational-style optimizer*

Pipelined Execution of Batch Programs

- *Better shuffle performance*
- *Scales to very large groups*

Active Library Development

- *Graphs / Machine Learning*
- *Streaming ML (coming)*

Native Iterations

- *Very fast Graph Processing*
- *Stateful Iterations for ML*

Technology inside Flink




- Technology inspired by compilers + MPP databases + distributed systems
- For ease of use, reliable performance, and scalability

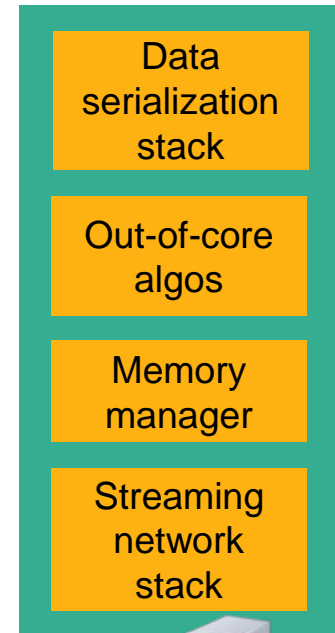
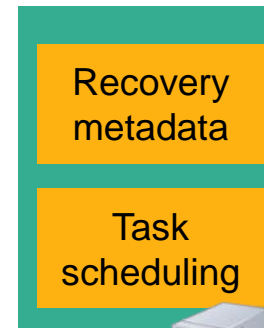
```
case class Path (from: Long, to: Long)
val tc = edges.iterate(10) {
  paths: DataSet[Path] =>
    val next = paths
      .join(edges)
      .where("to")
      .equalTo("from") {
        (path, edge) =>
          Path(path.from, edge.to)
      }
      .union(paths)
      .distinct()
    next
}
```





Pre-flight
(client)



Master



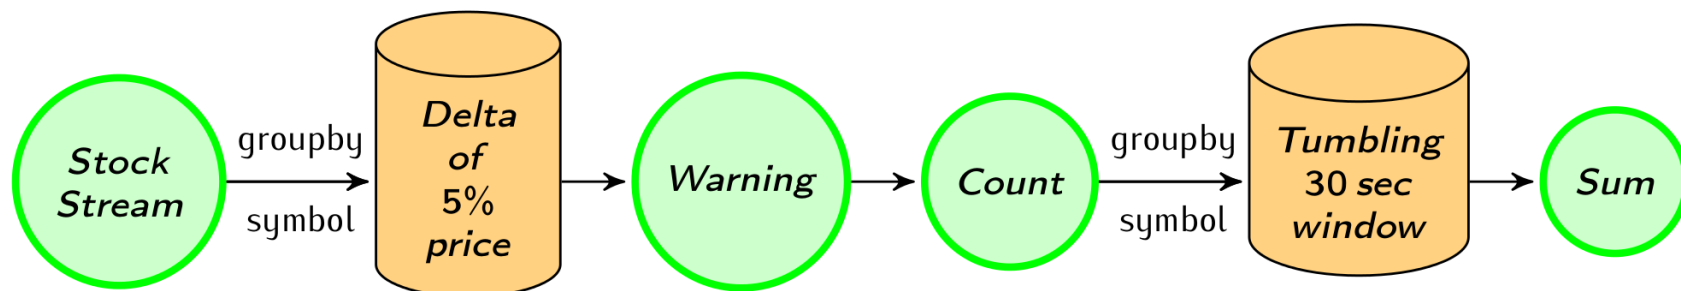
Workers





What can you do with Flink?

Streaming Data Analysis



```
case class Count(symbol: String, count: Int)
val defaultPrice = StockPrice("", 1000)

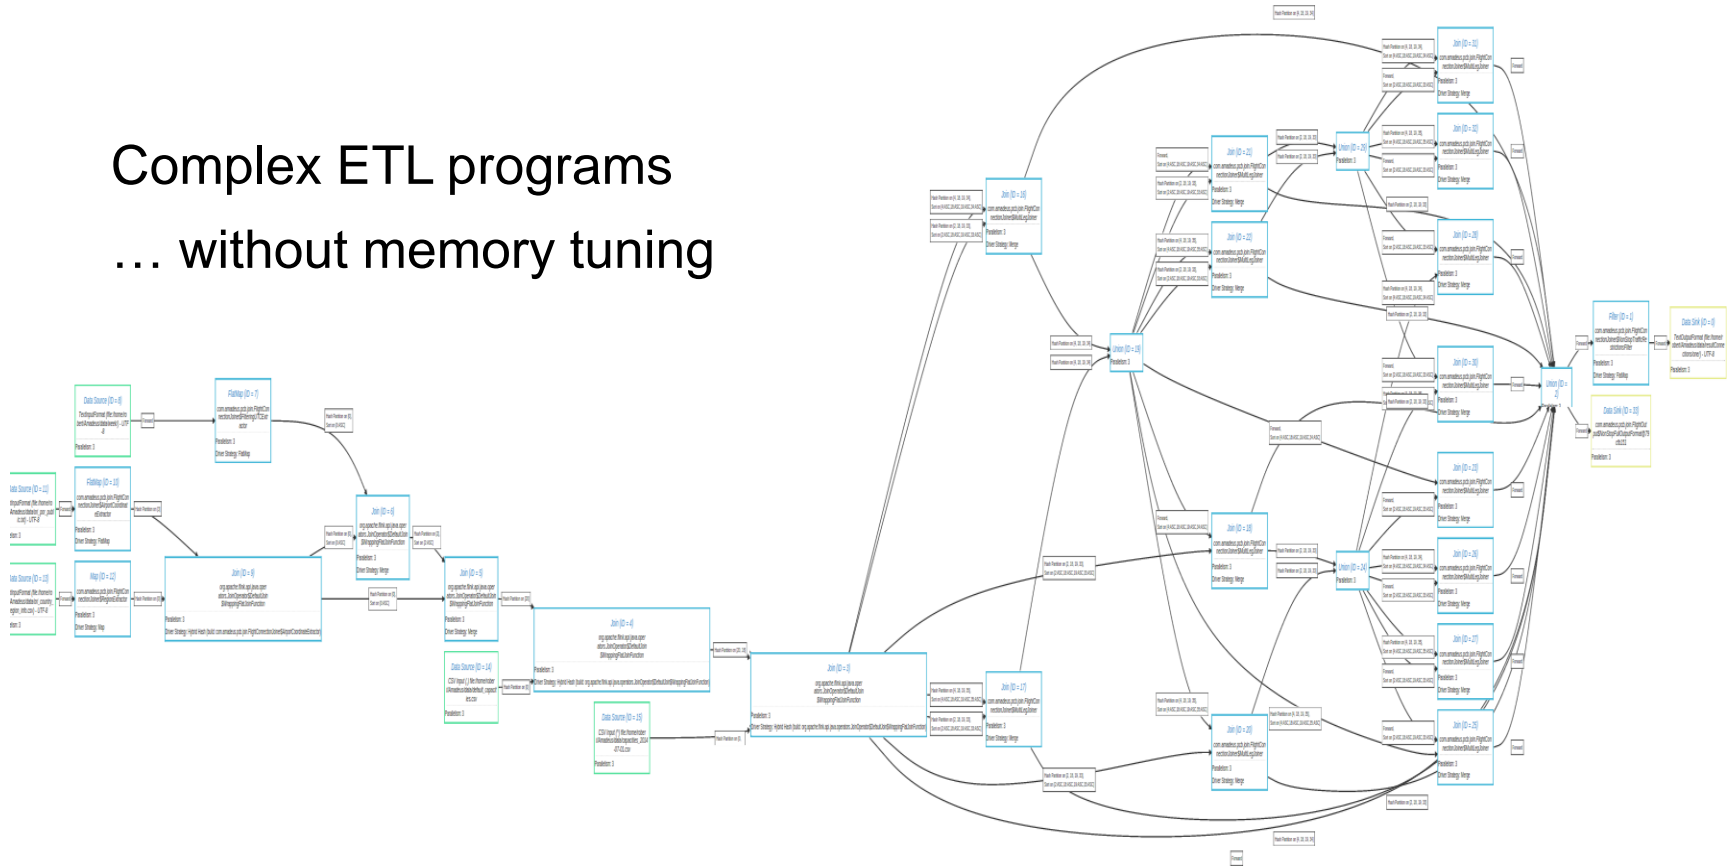
//Use delta policy to create price change warnings
val priceWarnings = stockStream.groupBy("symbol")
    .window(Delta.of(0.05, priceChange, defaultPrice))
    .mapWindow(sendWarning _)

//Count the number of warnings every half a minute
val warningsPerStock = priceWarnings.map(Count(_, 1))
    .groupBy("symbol")
    .window(Time.of(30, SECONDS))
    .sum("count")
```


Heavy Data Pipelines



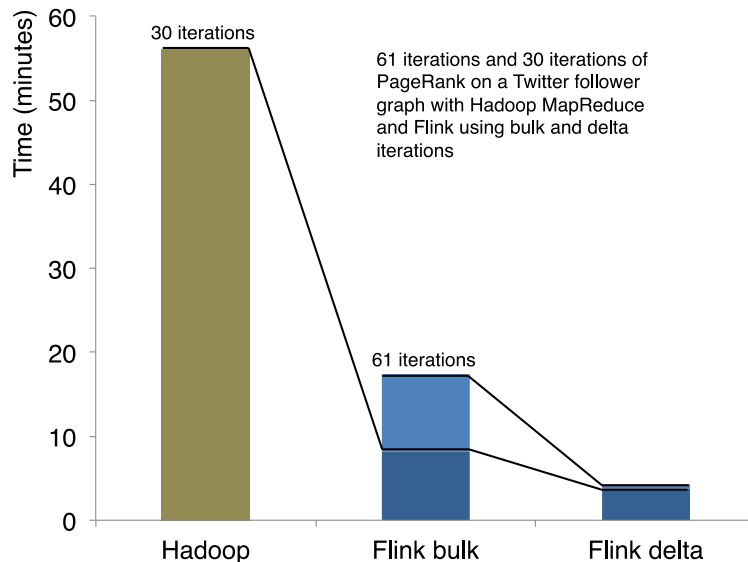

Complex ETL programs
... without memory tuning



Very fast graph analysis



... and mix and match
ETL-style and graph analysis
in one program



Performance competitive
with dedicated graph
analysis systems

More at: <http://data-artisans.com/data-analysis-with-flink.html>


Large-Scale Machine Learning



Factorizing a matrix with
28 billion ratings for
recommendations

		Item			
		W	X	Y	Z
User	A		4.5	2.0	
	B	4.0		3.5	
	C		5.0		2.0
	D		3.5	4.0	1.0
Rating Matrix					

		Item			
		W	X	Y	Z
User	A	1.2	0.8		
	B	1.4	0.9		
	C	1.5	1.0		
	D	1.2	0.8		
User Matrix		Item Matrix			



*(Scale of Netflix
or Spotify)*



How do you use Flink?

Example: WordCount



```
case class Word (word: String, frequency: Int)

val env = ExecutionEnvironment.getExecutionEnvironment()

val lines = env.readTextFile(...)

lines
    .flatMap {line => line.split(" ").map(word => Word(word,1))}
    .groupBy("word").sum("frequency")
    .print()

env.execute()
```

Flink has mirrored Java and Scala APIs that offer the same functionality, including by-name addressing.

Example: Window WordCount



```
case class Word (word: String, frequency: Int)

val env = StreamExecutionEnvironment.getExecutionEnvironment()

val lines = env.fromSocketStream(...)

lines
    .flatMap {line => line.split(" ").map(word => Word(word,1))}
    .window(Count.of(100)).every(Count.of(10))
    .groupBy("word").sum("frequency").print()

env.execute()
```

Flink API in a Nutshell



- map, flatMap, filter, groupBy, reduce, reduceGroup, aggregate, join, coGroup, cross, project, distinct, union, iterate, iterateDelta, ...
- All Hadoop input formats are supported
- API similar for data sets and data streams with slightly different operator semantics
- Window functions for data streams
- Counters, accumulators, and broadcast variables

Defining windows



- Trigger policy
 - When to trigger the computation on current window
- Eviction policy
 - When data points should leave the window
 - Defines window width/size
- E.g., count-based policy
 - evict when `#elements > n`
 - start a new window every `n`-th element
- Built-in: Count, Time, Delta policies

Table API



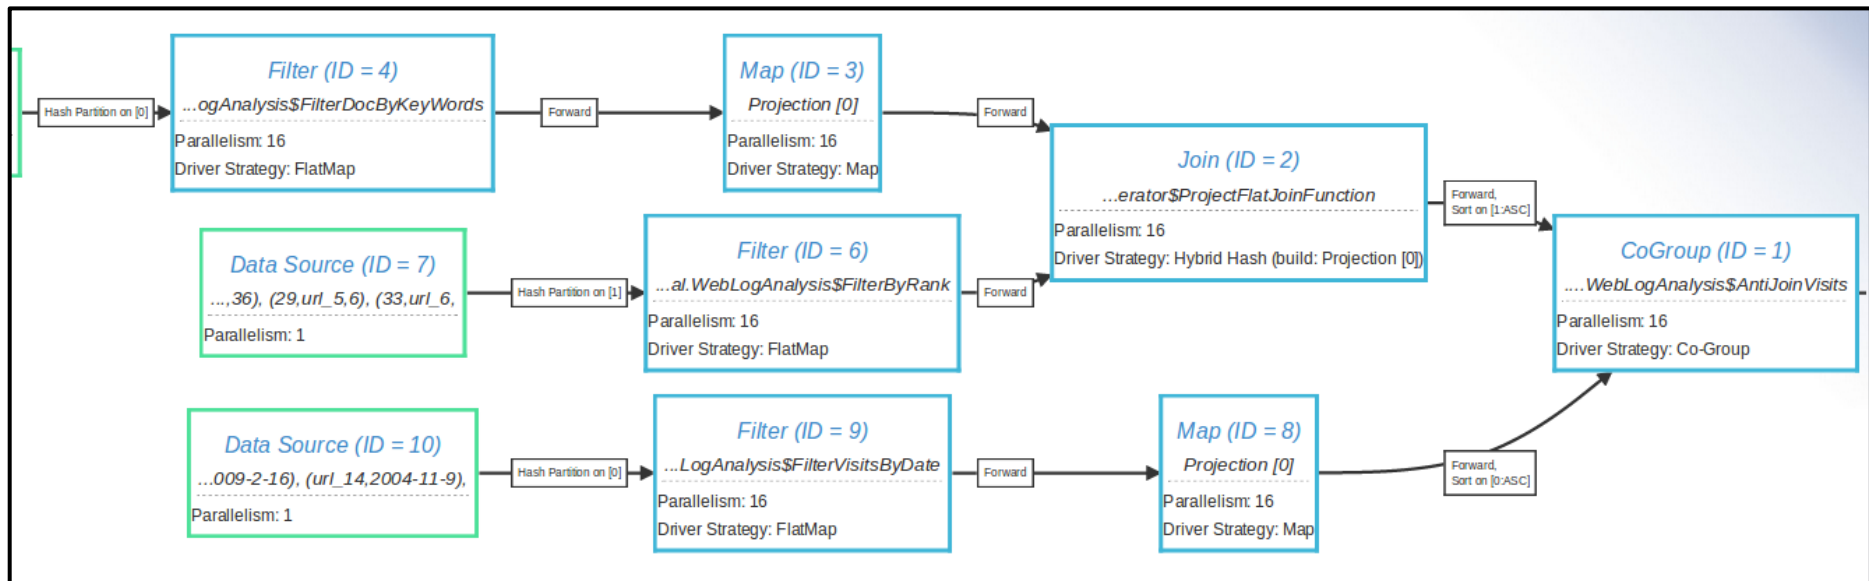
```
val customers = envreadCsvFile(...).as('id, 'mktSegment)
    .filter( 'mktSegment === "AUTOMOBILE" )

val orders = env.readCsvFile(...)
    .filter( o => dateFormat.parse(o.orderDate).before(date) )
    .as('orderId, 'custId, 'orderDate, 'shipPrio)

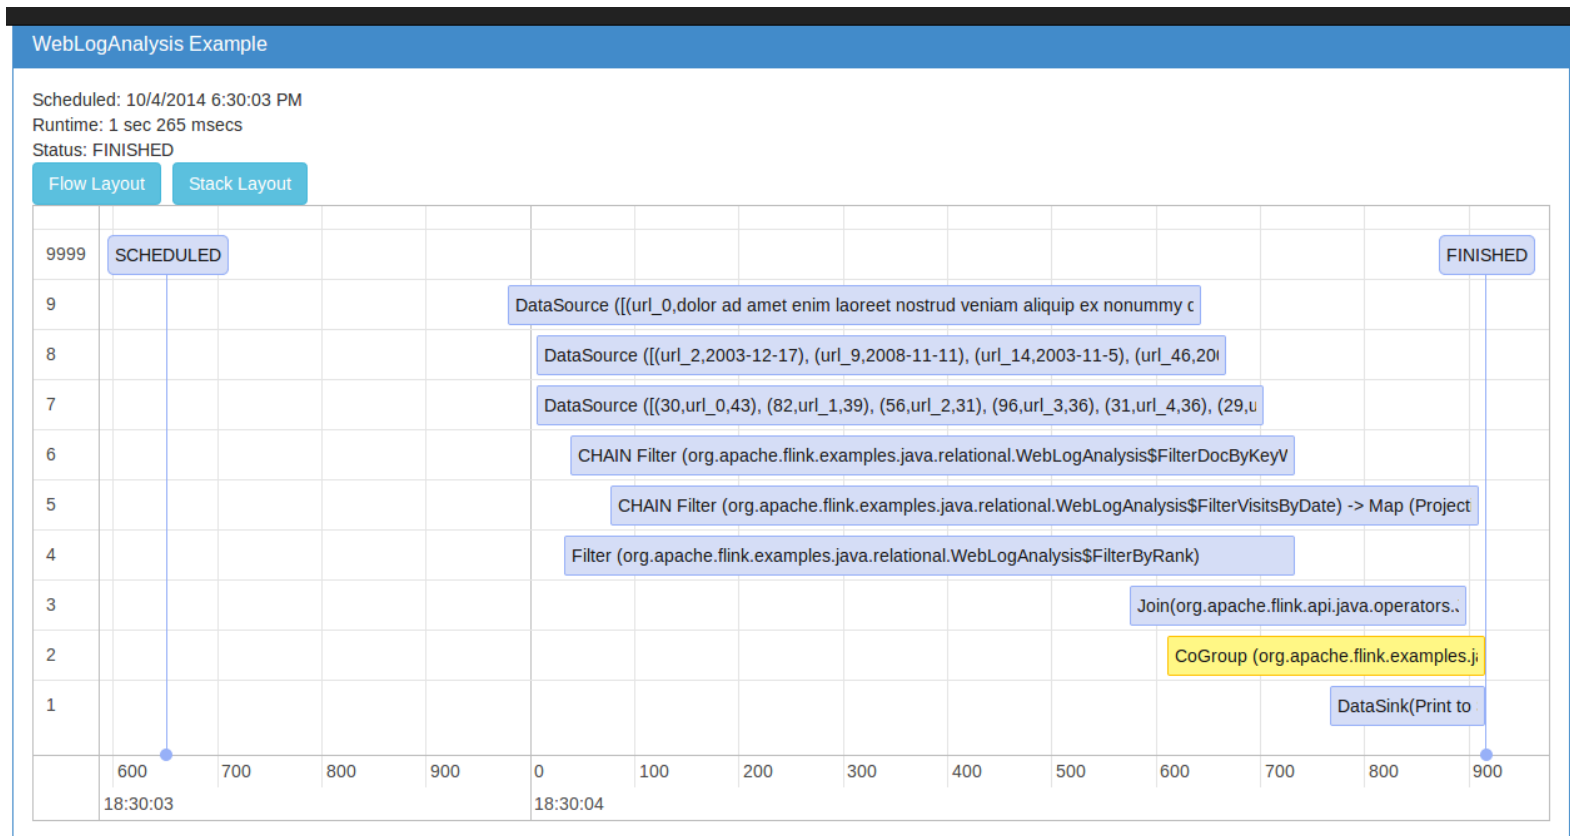
val items = orders
    .join(customers).where('custId === 'id)
    .join(lineitems).where('orderId === 'id)
    .select('orderId, 'orderDate, 'shipPrio,
        'extdPrice * (Literal(1.0f) - 'discount) as 'revenue)

val result = items
    .groupBy('orderId, 'orderDate, 'shipPrio)
    .select('orderId, 'revenue.sum, 'orderDate, 'shipPrio)
```



Visualization tools



Visualization tools



Visualization tools



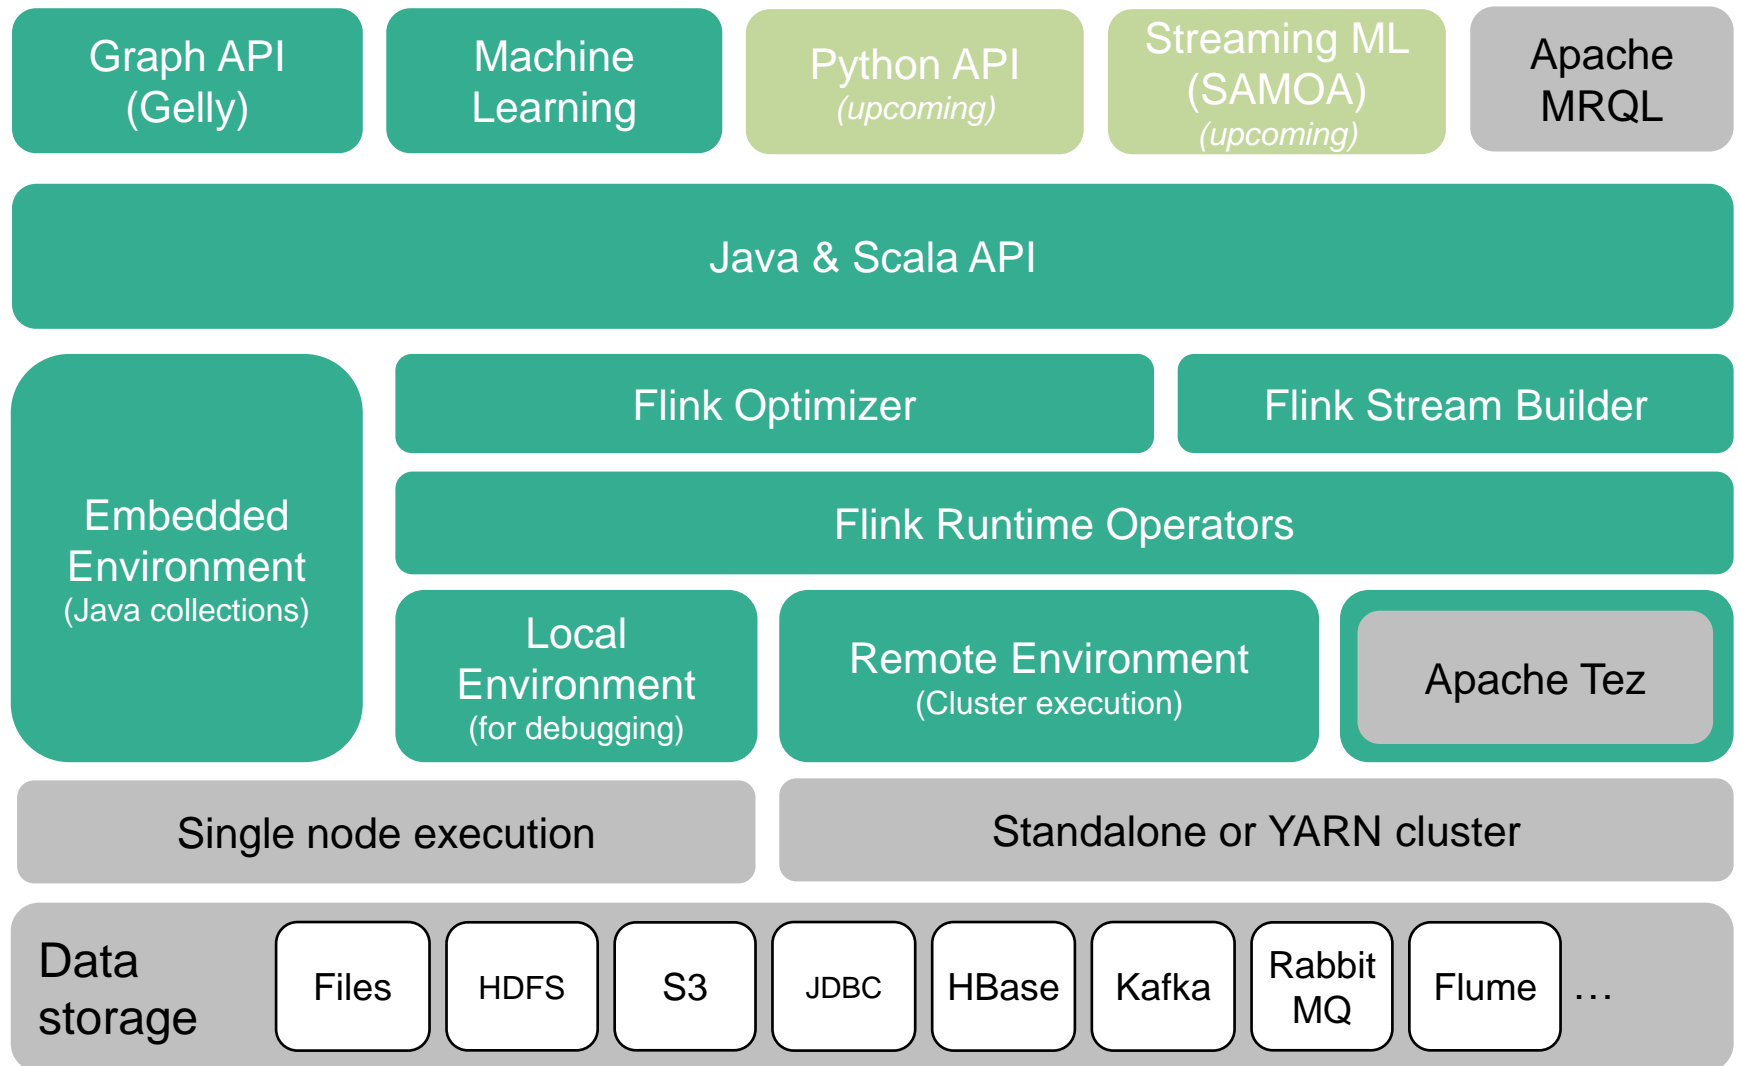
Flink Architecture

The case for Apache Flink



- Performance and ease of use
 - Exploits in-memory and pipelining, language-embedded logical APIs
- Unified batch and real streaming
 - Batch and Stream APIs on top of a streaming engine
- A runtime that "just works" without tuning
 - C++ style memory management inside the JVM
- Predictable and dependable execution
 - Bird's-eye view of what runs and how, and what failed and why

Flink stack





Evolution of Architectures



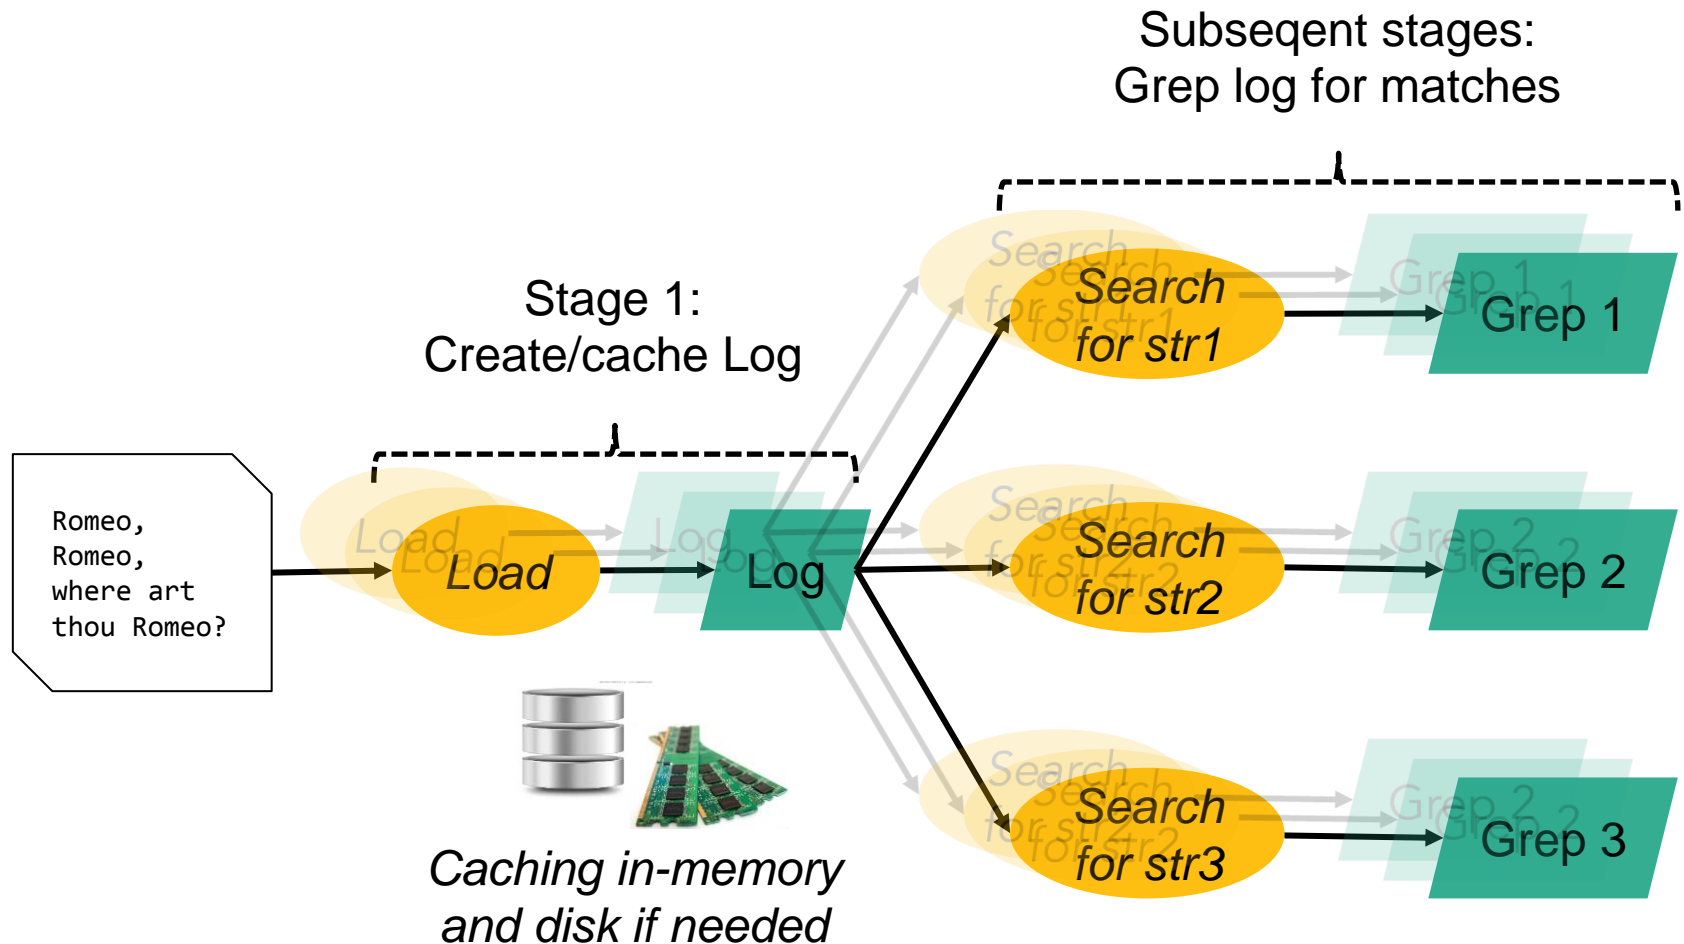
Operator-centric
(MapReduce / DAGs)

Dataset-centric
(RDDs)



**Operators and
DataSets**

Staged (batch) execution




Pipelined execution



Note: Log
DataSet is
never
“created”!

Stage 1:
Deploy and start operators



Pipelining in Flink




- Currently the default mode of operation
 - Much better performance in many cases – no need to materialize large data sets
 - Supports both batch and real-time streaming
- Currently evolving into a hybrid engine
 - Batch will use combination of blocking and pipelining
 - Streaming will use pipelining
 - Interactive will use blocking

Memory management

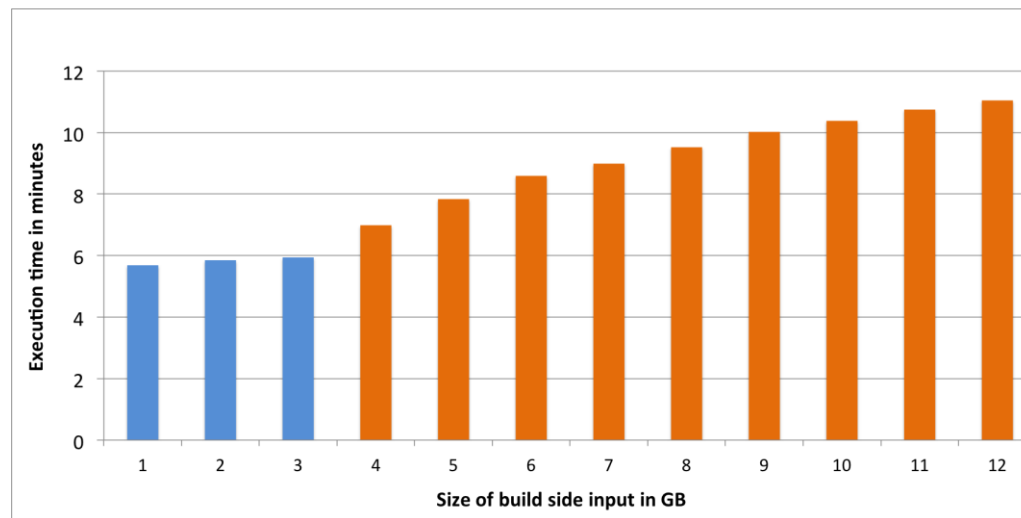
Memory management in Flink



Flink contains its own memory management stack. Memory is allocated, de-allocated, and used strictly using an internal buffer pool implementation. To do that, Flink contains its own type extraction and serialization components.



Smooth out-of-core performance

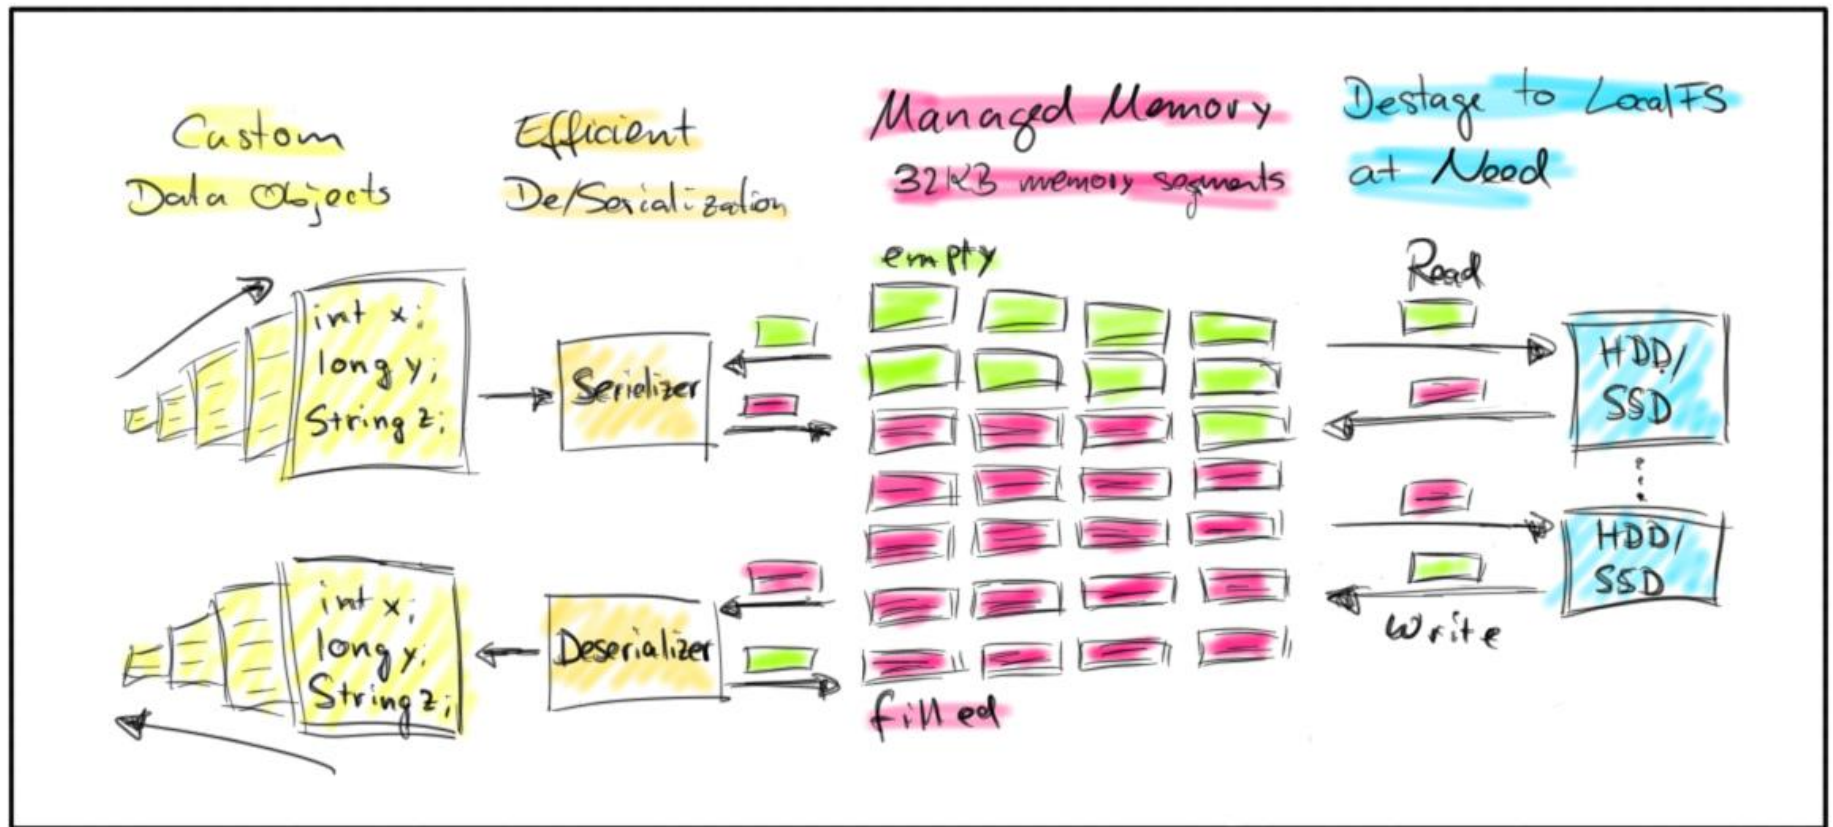


Single-core join of large Java objects beyond memory (4 GB)

Blue bars are in-memory, orange bars (partially) out-of-core

More at: <http://flink.apache.org/news/2015/03/13/peeking-into-Apache-Flinks-Engine-Room.html>

Paged Memory Management



Configuring Flink




- Per job
 - Parallelism
- System config
 - Total JVM heap size (-Xmx)
 - % of total JVM size for Flink runtime
 - Memory for network buffers (soon not needed)
- That's all you need. System will not throw an OOM exception to you.

Benefits of managed memory



- More reliable and stable performance (less GC effects, easy to go to disk)



Native iterative processing

Example: Transitive Closure



```
case class Path (from: Long, to: Long)

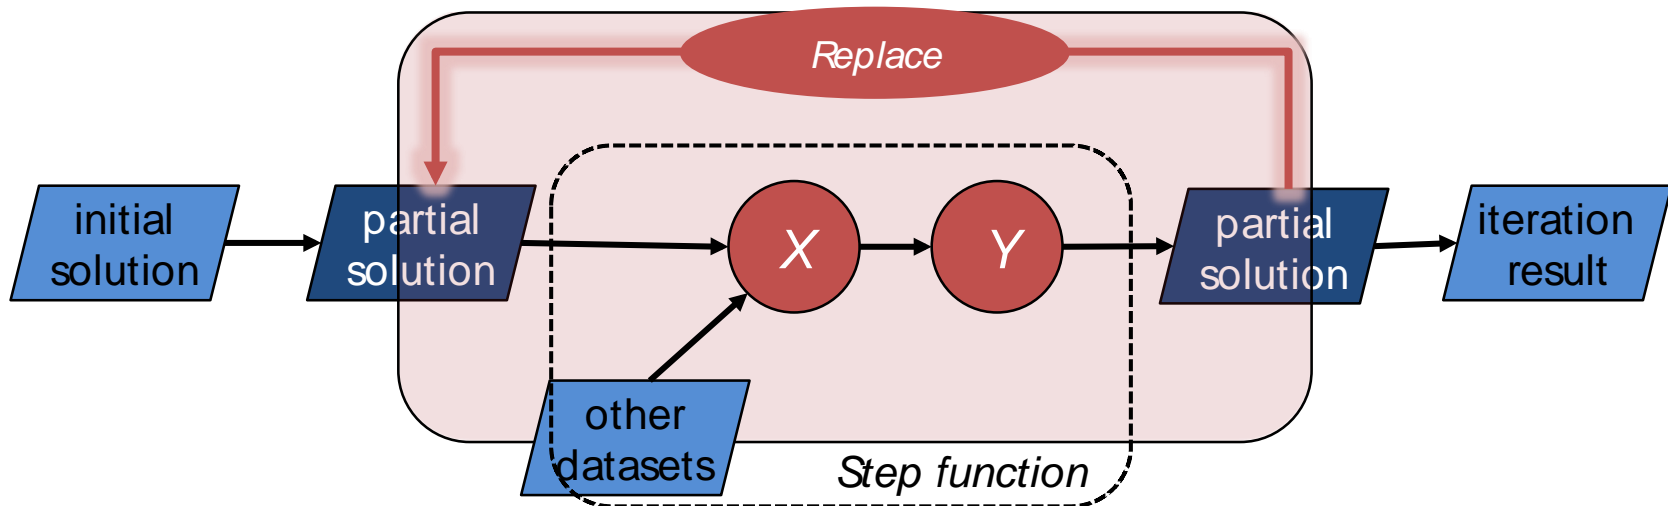
val env =
  ExecutionEnvironment.getExecutionEnvironment

val edges = ...


val tc = edges.iterate (10) { paths: DataSet[Path] =>
  val next = paths
    .join(edges).where("to").equalTo("from") {
      (path, edge) => Path(path.from, edge.to)
    }
    .union(paths).distinct()
  next
}

tc.print()
env.execute()
```


Iterate natively



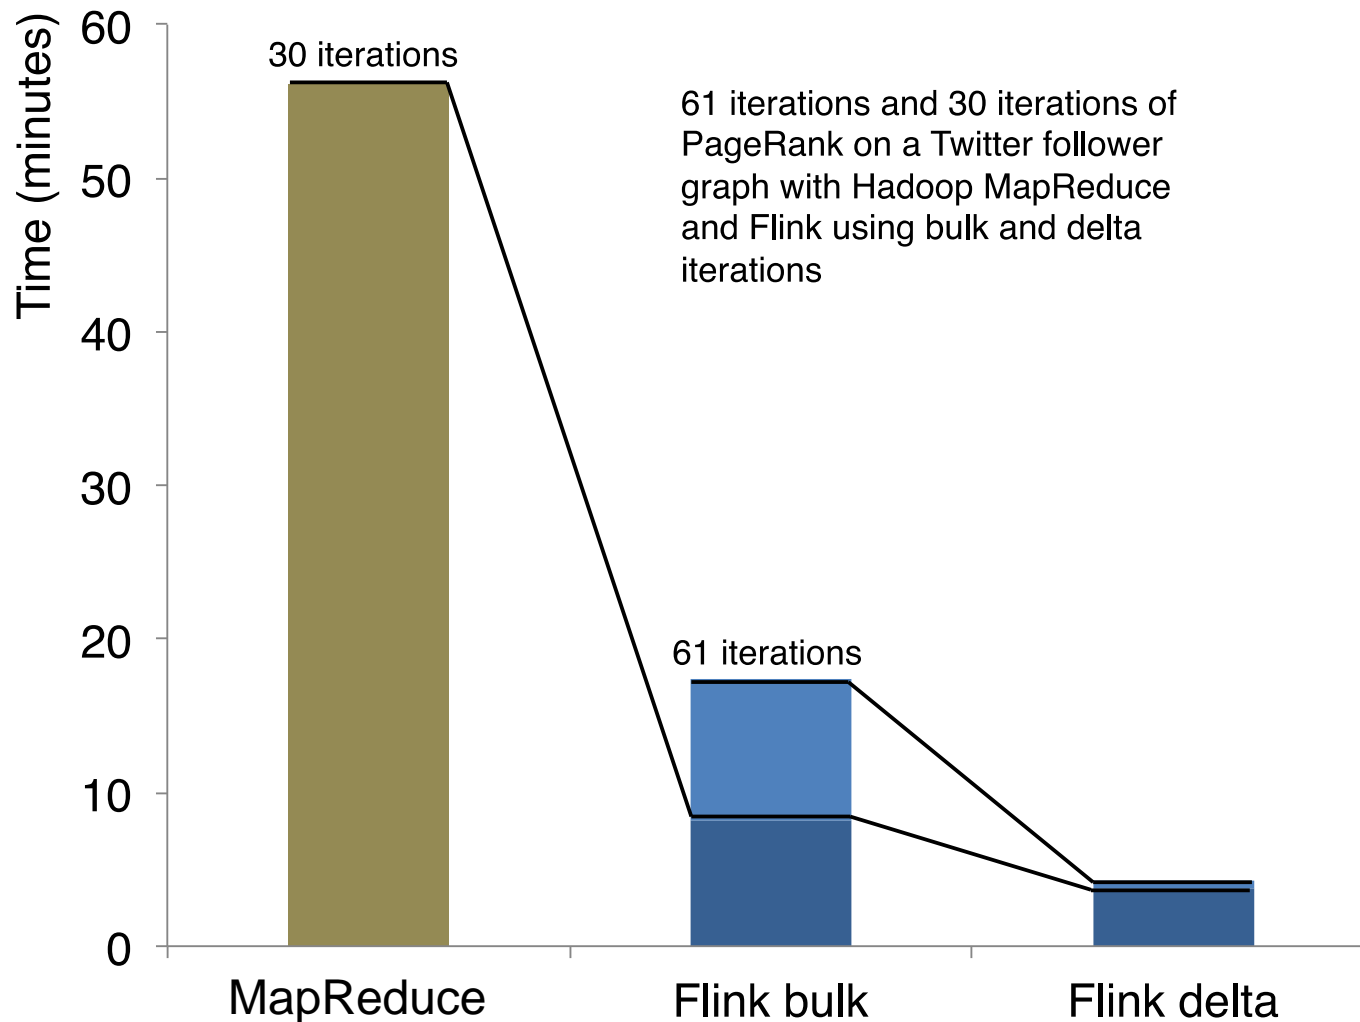
Iterate natively with deltas



Effect of delta iterations



Iteration performance



Program optimization

A simple program



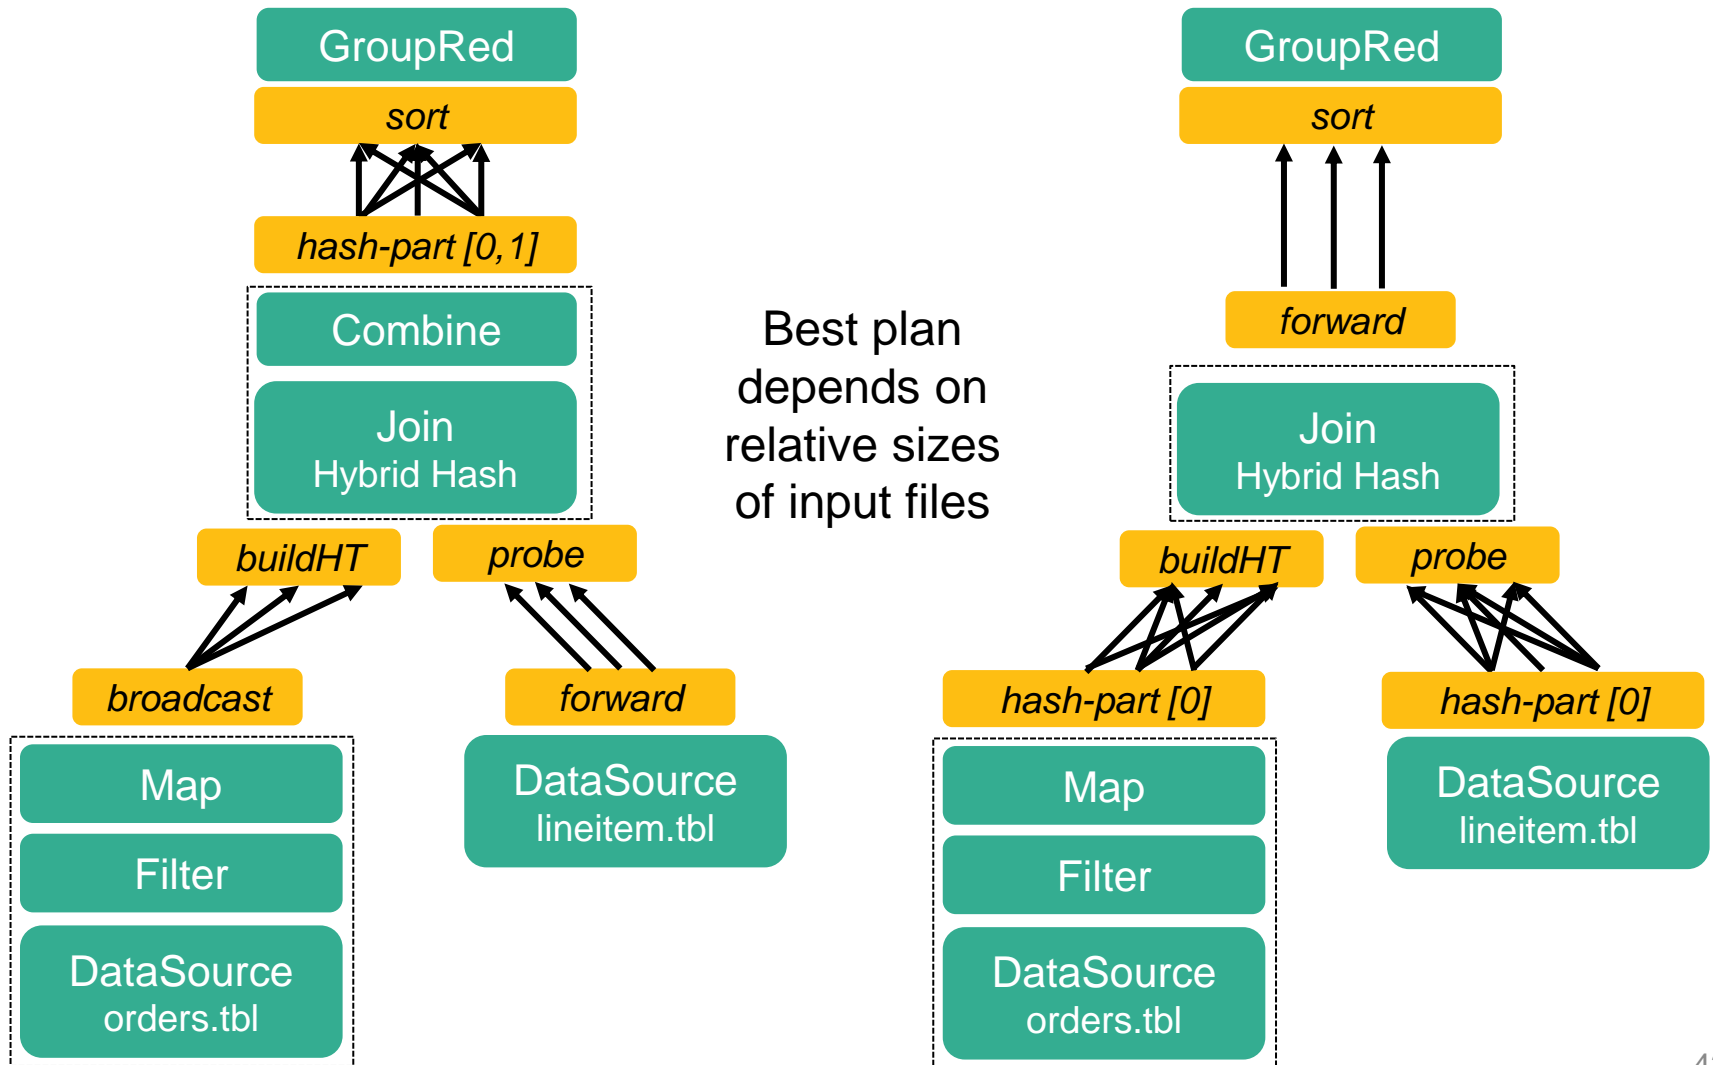
```
val orders = ...
val lineitems = ...

val filteredOrders = orders
  .filter(o => dateFormat.parse(l.shipDate).after(date))
  .filter(o => o.shipPrio > 2)

val lineitemsOfOrders = filteredOrders
  .join(lineitems)
  .where("orderId").equalTo("orderId")
  .apply((o,l) => new SelectedItem(o.orderDate, l.extdPrice))

val priceSums = lineitemsOfOrders
  .groupBy("orderDate").sum("l.extdPrice");
```

Two execution plans



Examples of optimization



- Task chaining
 - Coalesce map/filter/etc tasks
- Join optimizations
 - Broadcast/partition, build/probe side, hash or sort-merge
- Interesting properties
 - Re-use partitioning and sorting for later operations
- Automatic caching
 - E.g., for iterations

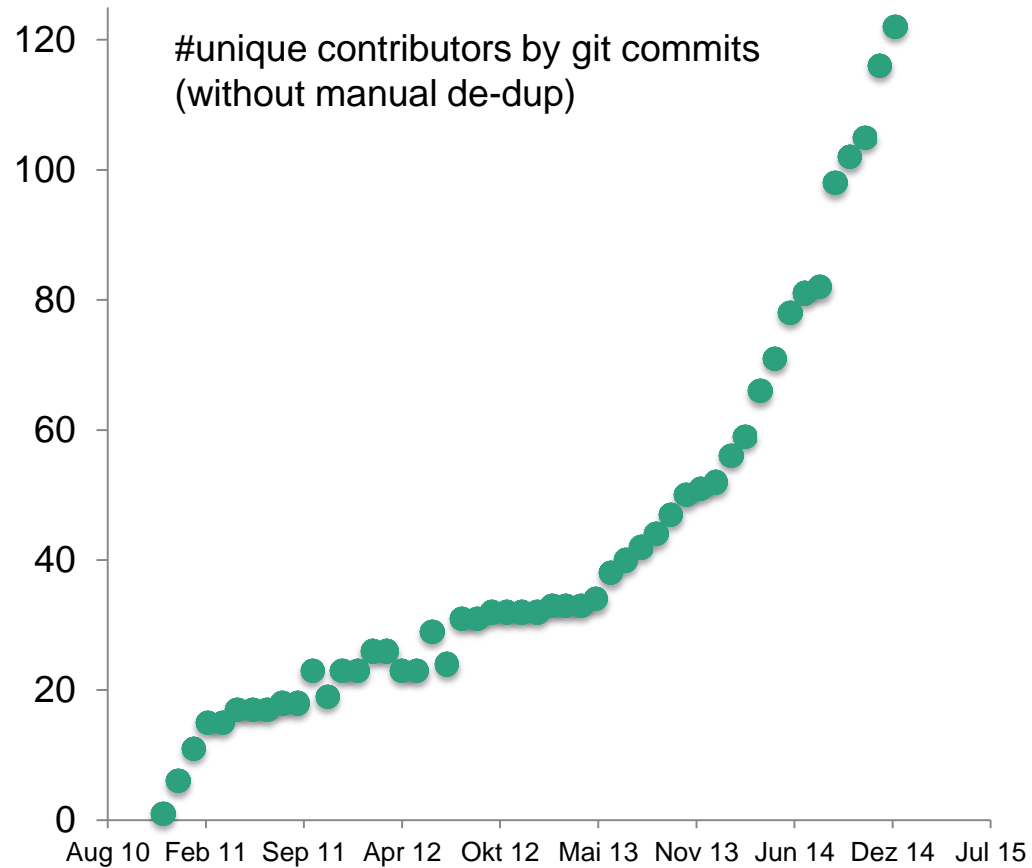
Closing

Flink Roadmap for 2015



- Exactly-once streaming with flexible state
- Support for Google Dataflow
- Batch Machine Learning library
- Streaming Machine Learning with SAMOA
- Graph library additions (more algorithms)
- Interactive programs and Zeppelin
- SQL on top of expression language
- and more...

Flink community





flink.apache.org
@ApacheFlink