

Real-time Stream Processing with Apache Flink



Marton Balassi – data Artisans
Gyula Fora - SICS

Flink committers

mbalassi@apache.org / gyfora@apache.org

Stream Processing



- **Data stream:** Infinite sequence of data arriving in a continuous fashion.
- **Stream processing:** Analyzing and acting on real-time streaming data, using continuous queries



Streaming landscape



Apache Storm

- True streaming, low latency - lower throughput
- Low level API (Bolts, Spouts) + Trident



Spark Streaming

- Stream processing on top of batch system, high throughput - higher latency
- Functional API (DStreams), restricted by batch runtime



Apache Samza

- True streaming built on top of Apache Kafka, state is first class citizen
- Slightly different stream notion, low level API



Apache Flink

- True streaming with adjustable latency-throughput trade-off
- Rich functional API exploiting streaming runtime; e.g. rich windowing semantics


Apache Storm



- True streaming, low latency - lower throughput
- Low level API (Bolts, Spouts) + Trident
- At-least-once processing guarantees

Issues

- Costly fault tolerance
- Serialization
- Low level API




Spark Streaming



- Stream processing emulated on a batch system
- High throughput - higher latency
- Functional API (DStreams)
- Exactly-once processing guarantees

Issues

- Restricted streaming semantics
- Windowing
- High latency




Apache Samza



- True streaming built on top of Apache Kafka
- Slightly different stream notion, low level API
- At-least-once processing guarantees with state

Issues


- High disk IO
- Low level API



Apache Flink



- True streaming with adjustable latency and throughput
- Rich functional API exploiting streaming runtime
- Flexible windowing semantics
- Exactly-once processing guarantees with (small) state



Issues

- Limited state size
- HA issue

Apache Flink



What is Flink




A "use-case complete" framework to unify
batch and stream processing



What is Flink



Real-time data
streams





Event logs

Kafka, RabbitMQ, ...

Historic data

HDFS, JDBC, ...

An engine that puts equal emphasis
to streaming and batch



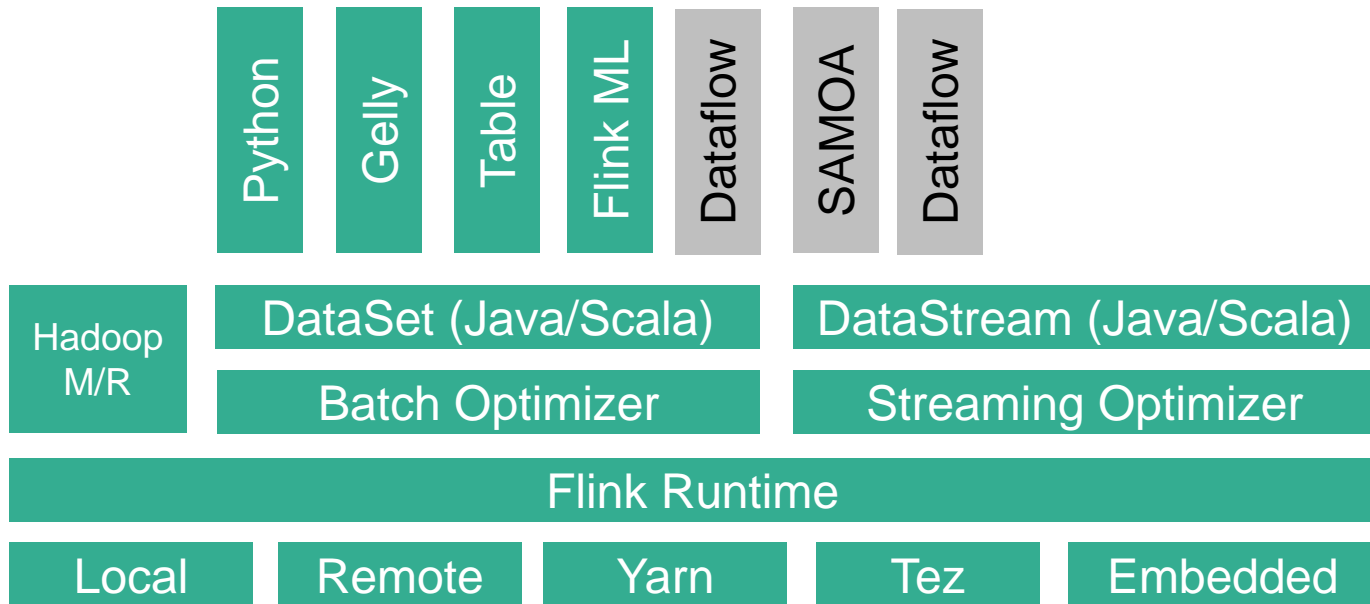
ETL, Graphs,
Machine Learning
Relational, ...

Low latency
windowing,
aggregations, ...

Flink stack



**current Flink master + few PRs*



Flink Streaming



Overview of the API



- Data stream sources

- File system
- Message queue connectors
- Arbitrary source functionality


- Stream transformations

- Basic transformations: *Map, Reduce, Filter, Aggregations...*
- Binary stream transformations: *CoMap, CoReduce...*
- Windowing semantics: *Policy based flexible windowing (Time, Count, Delta...)*
- Temporal binary stream operators: *Joins, Crosses...*
- Native support for iterations

- Data stream outputs

- For the details please refer to the programming guide:


- http://flink.apache.org/docs/latest/streaming_guide.html



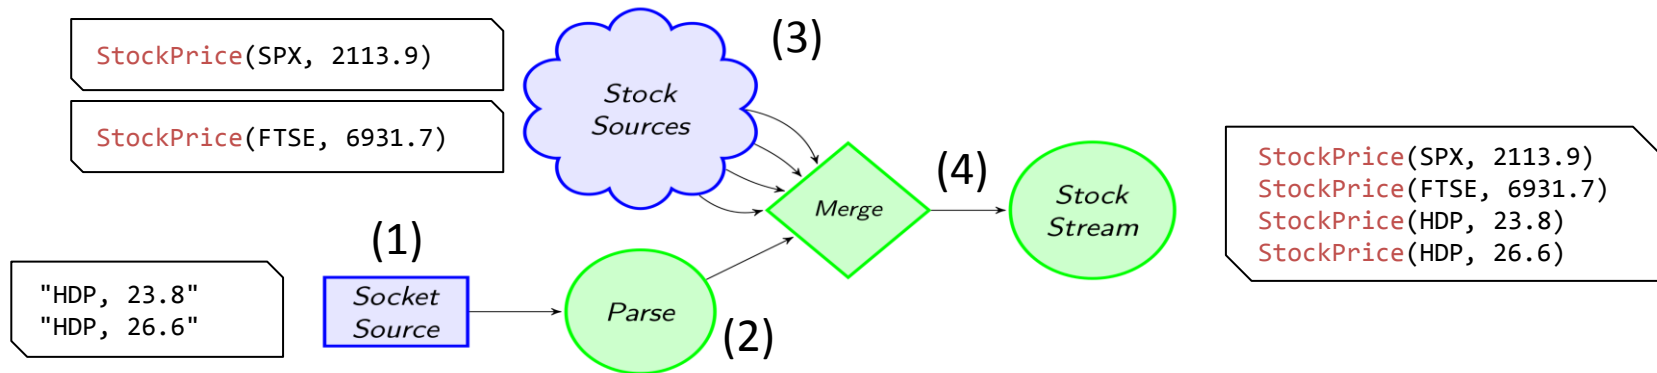
Use-case: Financial analytics



- Reading from multiple inputs
 - Merge stock data from various sources
- Window aggregations
 - Compute simple statistics over windows of data
- Data driven windows
 - Define arbitrary windowing semantics
- Combine with sentiment analysis
 - Enrich your analytics with social media feeds (Twitter)
- Streaming joins
 - Join multiple data streams
- Detailed explanation and source code on our blog
 - <http://flink.apache.org/news/2015/02/09/streaming-example.html>



Reading from multiple inputs



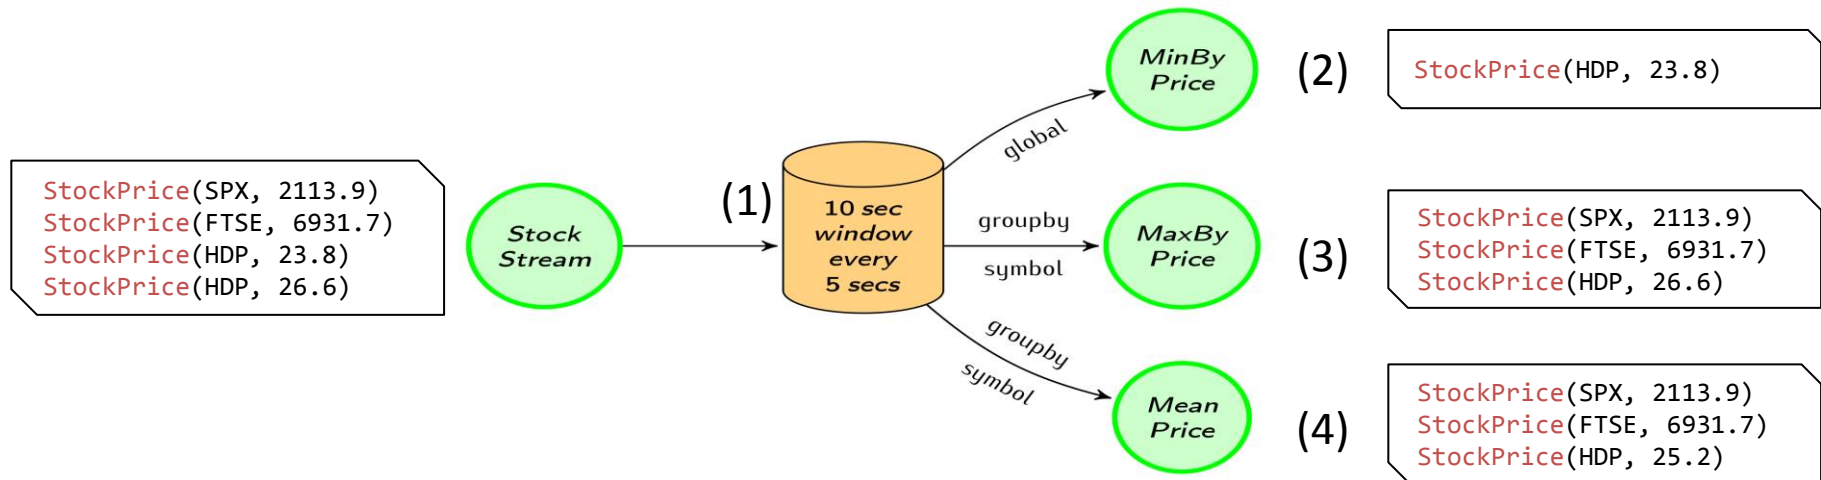
```
case class StockPrice(symbol : String, price : Double)
val env = StreamExecutionEnvironment.getExecutionEnvironment
```

```
(1) val socketStockStream = env.socketTextStream("localhost", 9999)
(2) { .map(x => { val split = x.split(",")
               StockPrice(split(0), split(1).toDouble) })
```

```
(3) { val SPX_Stream = env.addSource(generateStock("SPX")(10) _)
      val FTSE_Stream = env.addSource(generateStock("FTSE")(20) _)
```

```
(4) val stockStream = socketStockStream.merge(SPX_Stream, FTSE_Stream)
```


Window aggregations



```
val windowedStream = stockStream
(1) .window(Time.of(10, SECONDS)).every(Time.of(5, SECONDS))

(2) val lowest = windowedStream.minBy("price")
(3) val maxByStock = windowedStream.groupBy("symbol").maxBy("price")
(4) val rollingMean = windowedStream.groupBy("symbol").mapWindow(mean _)
```


Data-driven windows



```
case class Count(symbol : String, count : Int)
```

```
val priceWarnings = stockStream.groupBy("symbol")
```

```
(1) .window(Delta.of(0.05, priceChange, defaultPrice))
```

```
(2) .mapWindow(sendWarning _)
```

```
val warningsPerStock = priceWarnings.map(Count(_, 1)) .groupBy("symbol")
```


```
(3) .window(Time.of(30, SECONDS))
```

```
(4) .sum("count")
```

Combining with a Twitter stream



"hdp is on the rise!"
"I wish I bought more
YHOO and HDP stocks"




```
(1) val tweetStream = env.addSource(generateTweets _)
```

```
(2) {  
    (3) val mentionedSymbols = tweetStream.flatMap(tweet => tweet.split(" "))  
        .map(_.toUpperCase())  
        .filter(symbols.contains(_))  
}
```

```
val tweetsPerStock = mentionedSymbols.map(Count(_, 1)).groupBy("symbol")  
(4) .window(Time.of(30, SECONDS))  
    .sum("count")
```

Streaming joins



(1) {

```
val tweetsAndWarning = warningsPerStock.join(tweetsPerStock)
      .onWindow(30, SECONDS)
      .where("symbol")
      .equalTo("symbol"){ (c1, c2) => (c1.count, c2.count) }
```

(2) {

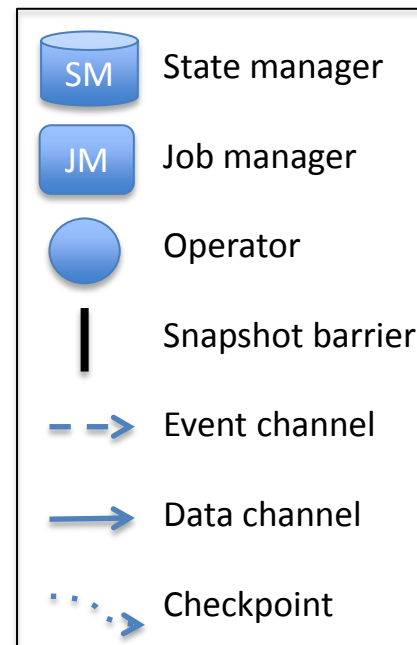

```
val rollingCorrelation = tweetsAndWarning
      .window(Time.of(30, SECONDS))
      .mapWindow(computeCorrelation _)
```

Fault tolerance



■ Exactly once semantics

- Asynchronous barrier snapshotting
- Checkpoint barriers streamed from the sources
- Operator state checkpointing + source backup
- Pluggable backend for state management



Performance



- Performance optimizations
 - Effective serialization due to strongly typed topologies
 - Operator chaining (thread sharing/no serialization)
 - Different automatic query optimizations
- Competitive performance
 - ~ 1.5m events / sec / core
 - As a comparison Storm promises ~ 1m tuples / sec / node

Roadmap



- Persistent, high-throughput state backend
- Job manager high availability
- Application libraries
 - General statistics over streams
 - Pattern matching
 - Machine learning pipelines library
 - Streaming graph processing library
- Integration with other frameworks
 - Zeppelin (Notebook)
 - SAMOA (Online ML)

Summary

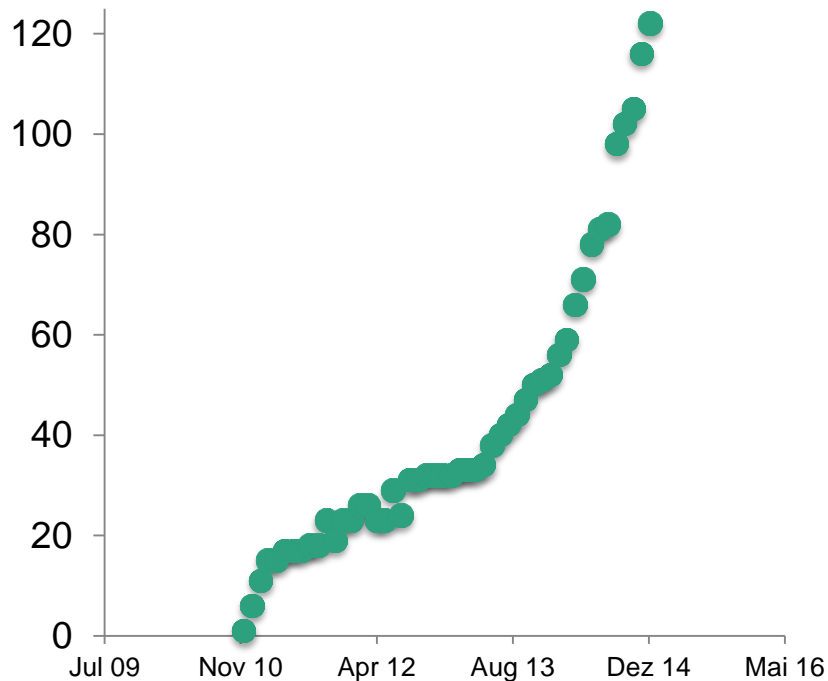


- Flink is a use-case complete framework to unify batch and stream processing
- True streaming runtime with high-level APIs
- Flexible, data-driven windowing semantics
- Competitive performance
- We are just getting started!

Flink Community



Unique git contributors



2015

Flink *Forward*

BERLIN 12/13 OCT



flink.apache.org
@ApacheFlink